

```

from sklearn.linear_model import LinearRegression
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score
from datetime import datetime
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error

```

```

from sklearn.linear_model import LinearRegression import matplotlib.pyplot as plt import numpy as np from
sklearn.model_selection import train_test_split from sklearn.metrics import mean_squared_error, r2_score from
datetime import datetime import pandas as pd stocks_data = pd.read_csv('stocks 1.csv') stocks_data

```

✓ EXPLORATORY DATA ANALYSIS

```

# @title EXPLORATORY DATA ANALYSIS
# Loading the data from a file
stocks_data = pd.read_csv('stocks 1.csv')

# Show the first few rows and basic information about the data
data_info = stocks_data.info() # Changed 'data' to 'stocks_data'
data_head = stocks_data.head() # Changed 'data' to 'stocks_data'

data_info, data_head

```



<class 'pandas.core.frame.DataFrame'>

RangeIndex: 248 entries, 0 to 247

Data columns (total 8 columns):

#	Column	Non-Null Count	Dtype
0	Ticker	248 non-null	object
1	Date	248 non-null	object
2	Open	248 non-null	float64
3	High	248 non-null	float64
4	Low	248 non-null	float64
5	Close	248 non-null	float64
6	Adj Close	248 non-null	float64
7	Volume	248 non-null	int64

dtypes: float64(5), int64(1), object(2)

memory usage: 15.6+ KB

(None,

	Ticker	Date	Open	High	Low	Close	\
0	AAPL	07-02-2023	150.639999	155.229996	150.639999	154.649994	
1	AAPL	08-02-2023	153.880005	154.580002	151.169998	151.919998	
2	AAPL	09-02-2023	153.779999	154.330002	150.419998	150.869995	
3	AAPL	10-02-2023	149.460007	151.339996	149.220001	151.009995	
4	AAPL	13-02-2023	150.949997	154.259995	150.919998	153.850006	

	Adj Close	Volume
0	154.414230	83322600
1	151.688400	64120100
2	150.639999	56007100
3	151.009995	57450700
4	153.850006	62199000

```
stocks_data.Ticker.value_counts()
```




	count
Ticker	
AAPL	62
MSFT	62
NFLX	62
GOOG	62



```
descriptive_stats = stocks_data.groupby('Ticker')
descriptive_stats['Close'].describe()
```



	count	mean	std	min	25%	50%	75%	max
Ticker								
AAPL	62.0	158.240645	7.360485	145.309998	152.077499	158.055000	165.162506	173.570007
GOOG	62.0	100.631532	6.279464	89.349998	94.702501	102.759998	105.962503	109.459999
MSFT	62.0	275.039839	17.676231	246.270004	258.742500	275.810013	287.217506	310.649994
NFLX	62.0	327.614677	18.554410	292.760010	315.672403	325.600006	338.800004	366.820087



```
# Sort the data by the 'Volume' column in descending order
sorted_data_volume = stocks_data.sort_values(by='Volume', ascending=False) # Changed 'data' to 'stocks_data'

# Display the top 5 rows with the highest trading volume
top_5_volume = sorted_data_volume[['Ticker', 'Date', 'Volume']].head(5)

# Print the top 5 dates with the most stocks sold
print("Top 5 dates with the most stocks sold:")
print(top_5_volume)
```



Top 5 dates with the most stocks sold:

	Ticker	Date	Volume
61	AAPL	05-05-2023	113316400
27	AAPL	17-03-2023	98944600
188	GOOG	09-02-2023	97798600
18	AAPL	06-03-2023	87558000
23	AAPL	13-03-2023	84457100

```
# Calculate the profit for each stock (difference between closing price and opening price)
stocks_data['Profit'] = stocks_data['Close'] - stocks_data['Open'] # Changed 'data' to 'stocks_data'

# Group by the stock ticker and sum the total profit for each stock
stock_profit = stocks_data.groupby('Ticker')['Profit'].sum().sort_values(ascending=False) # Changed 'data' to 'stocks_data'

# Display the profit for each stock ticker
print("Total profit for each stock ticker:")
print(stock_profit)
```



Total profit for each stock ticker:

Ticker	
AAPL	28.569946
MSFT	18.839966
GOOG	15.476021
NFLX	-30.749908
Name: Profit, dtype: float64	

```
# Sort the data by the 'High' price in descending order
sorted_data = stocks_data.sort_values(by='High', ascending=False) # Changed 'data' to 'stocks_data'

# Display the top 5 rows with the highest stock prices
top_5_high_prices = sorted_data[['Ticker', 'Date', 'High']].head(5)
```

```
# Print the top 5 high prices
print("Top 5 highest stock prices:")
print(top_5_high_prices)
```

↗ Top 5 highest stock prices:

	Ticker	Date	High
126	NFLX	09-02-2023	373.829987
125	NFLX	08-02-2023	368.190002
124	NFLX	07-02-2023	364.179993
129	NFLX	14-02-2023	363.750000
130	NFLX	15-02-2023	362.880005

```
# Sort the data by the 'Low' price in ascending order
sorted_data_low = stocks_data.sort_values(by='Low', ascending=True) # Changed 'data' to 'stock_data'

# Display the top 5 rows with the lowest stock prices
top_5_low_prices = sorted_data_low[['Ticker', 'Date', 'Low']].head(5)

# Print the top 5 least low stock prices
print("Top 5 least low stock prices:")
print(top_5_low_prices)
```

↗ Top 5 least low stock prices:

	Ticker	Date	Low
198	GGOG	24-02-2023	88.860001
200	GGOG	28-02-2023	89.519997
199	GGOG	27-02-2023	89.610001
202	GGOG	02-03-2023	89.769997
201	GGOG	01-03-2023	89.849998

✓ LINEAR REGRESSION

```
# @title LINEAR REGRESSION
# Convert the 'Date' column to datetime format
stocks_data['Date'] = pd.to_datetime(stocks_data['Date'], format='%d-%m-%Y')
```

```
# Convert dates to ordinal (numeric) values for regression
stocks_data['Date_ordinal'] = stocks_data['Date'].apply(lambda date: date.toordinal())
```

```
# Extract the features (Date) and target variable (Adj Close)
X = stocks_data[['Date_ordinal']]
y = stocks_data['Adj Close']
```

```
# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# Create and fit the linear regression model
model = LinearRegression()
model.fit(X_train, y_train)
```

↗

▼ LinearRegression ⓘ ?
 LinearRegression()


```
# Predict the target values for test data
y_pred = model.predict(X_test)
```

```
# Calculate performance metrics
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

# Plot the actual vs predicted values and the best-fit line
plt.figure(figsize=(10, 6))
plt.scatter(X_test, y_test, color='blue', label='Actual Prices')
plt.plot(X_test, y_pred, color='red', label='Best Fit Line')
```

```
plt.title('Linear Regression on Stock Prices (Adj Close)')
plt.xlabel('Date (Ordinal)')
plt.ylabel('Adjusted Close Price')
plt.legend()
plt.grid(True)
plt.show()

mse, r2
```



(8310.036359255038, -0.019042937503605195)

```
# Calculate additional regression metrics: Mean Absolute Error (MAE) and Mean Absolute Percentage Error (MAPE)
from sklearn.metrics import mean_absolute_error

# Mean Absolute Error (MAE)
mae = mean_absolute_error(y_test, y_pred)

# Mean Absolute Percentage Error (MAPE)
mape = np.mean(np.abs((y_test - y_pred) / y_test)) * 100

mae, mape
```



(87.246631804702, 56.08417736043175)

✓ LOGISTIC REGRESSION

```
# @title LOGISTIC REGRESSION
# Create the binary target variable
stocks_data['Price_Up'] = (stocks_data['Adj Close'].shift(-1) > stocks_data['Adj Close']).astype(int)

# Drop the last row as it won't have a valid comparison
stocks_data = stocks_data.dropna()

# Define the feature set and target variable
X = stocks_data[['Open', 'High', 'Low', 'Close', 'Volume']] # Using numerical features
y = stocks_data['Price_Up']
```

```
# Split the data into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# Initialize and fit the logistic regression model
model = LogisticRegression()
model.fit(X_train, y_train)
```



LogisticRegression ⓘ ?

LogisticRegression()

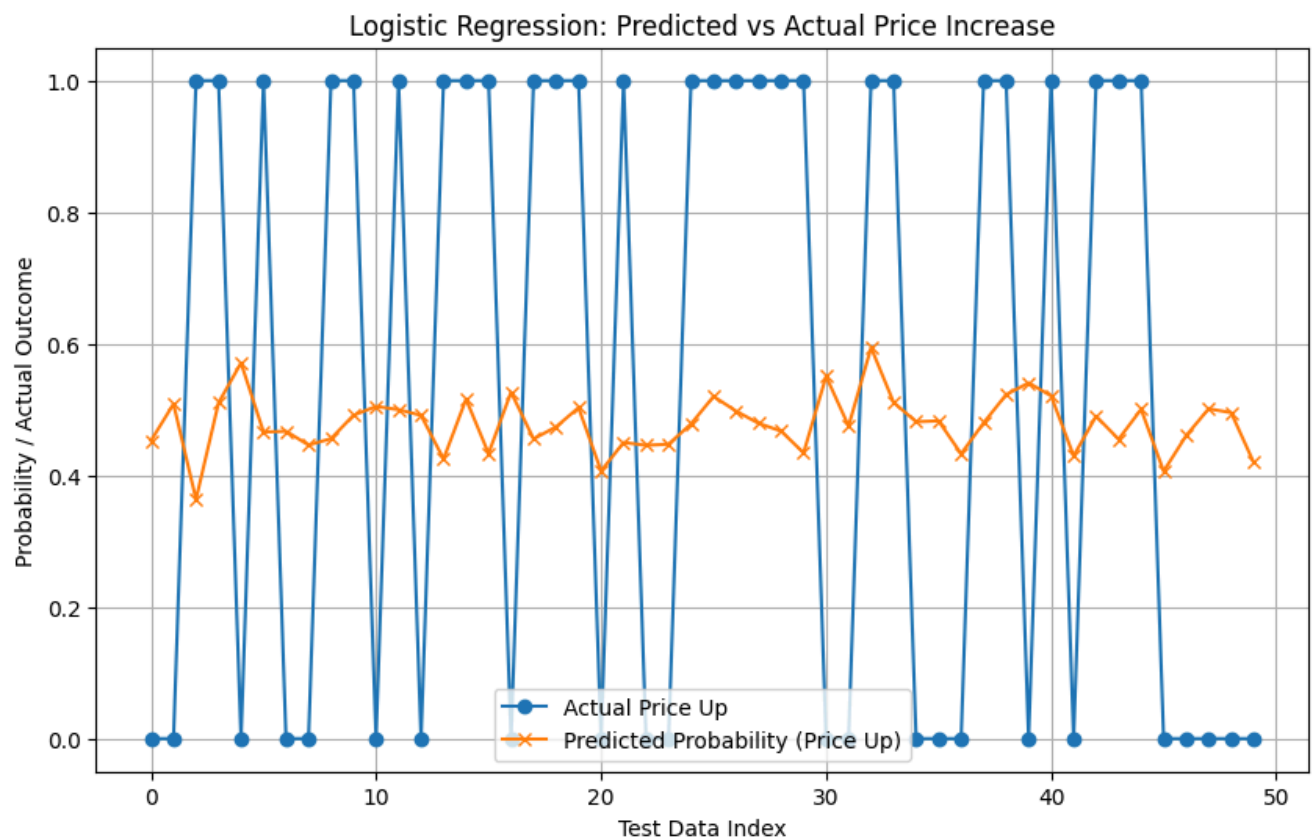
```
# Make predictions on the test set
y_pred = model.predict(X_test)
```

```
# Initialize and fit the logistic regression model
model = LogisticRegression()
model.fit(X_train, y_train)
```

```
# Get the predicted probabilities for class 1 (Price Up)
y_prob = model.predict_proba(X_test)[:, 1]
```

```
# Plot the predicted probabilities and actual values
plt.figure(figsize=(10, 6))
plt.plot(y_test.reset_index(drop=True), label='Actual Price Up', marker='o')
plt.plot(y_prob, label='Predicted Probability (Price Up)', marker='x')
```

```
plt.title('Logistic Regression: Predicted vs Actual Price Increase')
plt.xlabel('Test Data Index')
plt.ylabel('Probability / Actual Outcome')
plt.legend()
plt.grid(True)
plt.show()
```



```
# Generate the classification report to evaluate performance
classification_report_result = classification_report(y_test, y_pred)

classification_report_result
```

		precision	recall	f1-score	support		0	0.48	0.70	0.57
23	1	0.59	0.37	0.45	27	accuracy				0.
52	50	macro avg	0.54	0.53	0.51	50	weighted avg	0.54	0.52	


```
#time series analysis

#time series analysis
import plotly.express as px
from plotly.subplots import make_subplots
import plotly.graph_objects as go
import pandas as pd

# Convert the existing 'Date' column to datetime
# Check for typos or case sensitivity issues in the column name
stocks_data['Date'] = pd.to_datetime(stocks_data['Date'], format='%d-%m-%Y') # Changed 'Date' to 'date'
pivot_data = stocks_data.pivot(index='Date', columns='Ticker', values='Close')
fig = make_subplots(rows=1, cols=1)
fig.add_trace(go.Scatter(x=pivot_data.index, y=pivot_data['AAPL'], name='AAPL'))
fig.add_trace(go.Scatter(x=pivot_data.index, y=pivot_data['GOOG'], name='GOOG'))
fig.add_trace(go.Scatter(x=pivot_data.index, y=pivot_data['NFLX'], name='NFLX'))
fig.add_trace(go.Scatter(x=pivot_data.index, y=pivot_data['MSFT'], name='MSFT'))
fig.update_layout(
    title_text="Time Series of Closing Prices",
    xaxis_title='Date',
    yaxis_title='Closing Price',
    legend_title='Ticker',
    showlegend=True
)

fig.show()
```

Time Series of Closing Prices



✓ VOLATILITY ANALYSIS

```
# @title VOLATILITY ANALYSIS
#volatility analysis
volatility = pivot_data.std()
fig = px.bar(
    volatility,
    x=volatility.index,
    y=volatility.values,
```

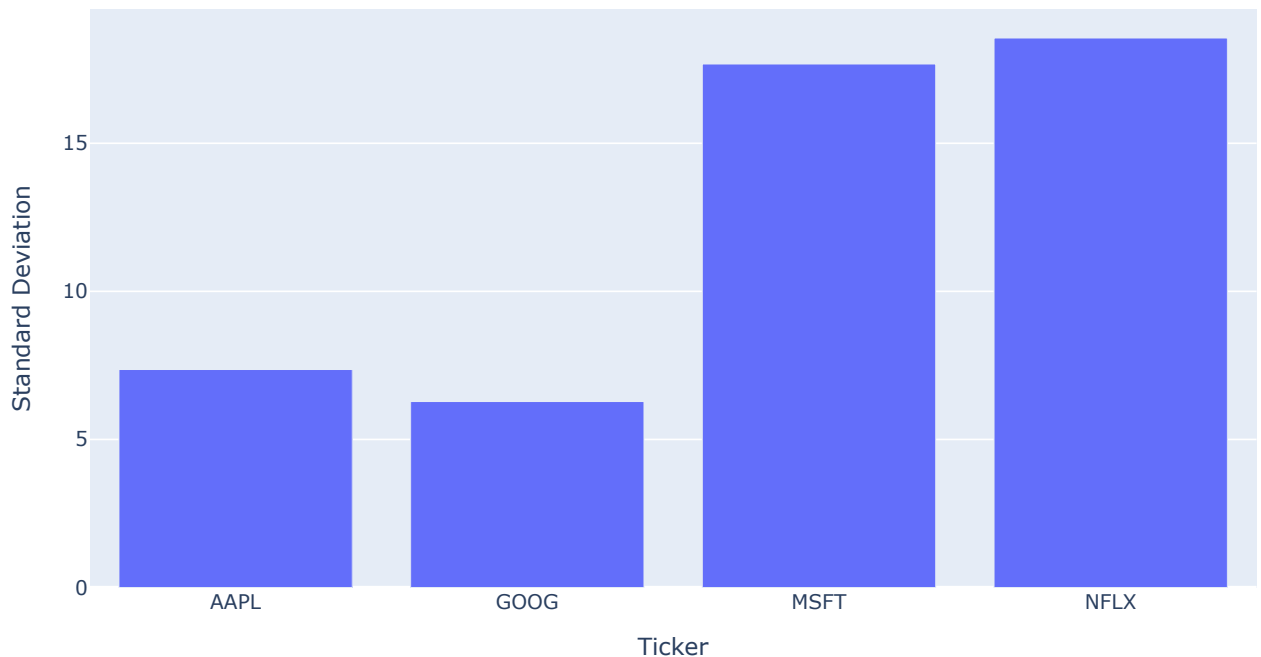
```

labels={
    'y':'Standard Deviation',
    'x':'Ticker'
},
title='Volatility of Closing Prices (Standard Deviation)'
)
fig.show()

```



Volatility of Closing Prices (Standard Deviation)



✖ CORRELAATION ANALYSIS


```

# @title CORRELAATION ANALYSIS
#correlation analysis
correlation_matrix = pivot_data.corr()
fig = go.Figure(
    data=go.Heatmap(
        z=correlation_matrix,
        x=correlation_matrix.columns,
        y=correlation_matrix.columns,
        colorscale='blues',
        colorbar=dict(title='correlation'),
        text=correlation_matrix.round(2).values,
        texttemplate="%{text}"
    )
)
fig.update_layout(
    title='Correlation Matrix of Closing Prices',
    xaxis_title="Ticker",
    yaxis_title="Ticker",
)
fig.show()

```



Correlation Matrix of Closing Prices



COMPARITIVE ANALYSIS

```
# @title COMPARITIVE ANALYSIS
#comparitive analysis
# Calculating the percentage change in closing prices
percentage_change = ((pivot_data.iloc[-1] - pivot_data.iloc[0]) / pivot_data.iloc[0]) * 100
fig = px.bar(
    percentage_change,
    x=percentage_change.index,
    y=percentage_change.values,
    labels={'y': 'Percentage Change (%)', 'x': 'Ticker'},
    title='Percentage Change in Closing Prices'
)
fig.show()
```




Percentage Change in Closing Prices



▼ RISK vs RETURN

```
# @title RISK vs RETURN
#daily risks vs return risks
daily_returns = pivot_data.pct_change().dropna()
avg_daily_return = daily_returns.mean()
risk = daily_returns.std()


risk_return_df = pd.DataFrame({'Risk':risk,'Average Daily Return':avg_daily_return})

fig = go.Figure()

fig.add_trace(
    go.Scatter(
        x=risk_return_df["Risk"],
        y=risk_return_df['Average Daily Return'],
        mode="markers+text",
        text=risk_return_df.index,
        textposition="top center",
        marker=dict(size=10)
    )
)
fig.update_layout(
    title='Risk vs. Return Analysis',
    xaxis_title='Risk (Standard Deviation)',
    yaxis_title='Average Daily Return',
    showlegend=False
)
fig.show()
```



Risk vs. Return Analysis



RANDOM FOREST REGRESSION

```
# @title RANDOM FOREST REGRESSION
# Let's assume 'Date' is a column in your dataset and 'Close' is the target feature for prediction
# You may need to adjust the column names based on your dataset

# Convert the 'Date' column to datetime (if present)
if 'Date' in stocks_data.columns:
    stocks_data['Date'] = pd.to_datetime(stocks_data['Date'])
    stocks_data.set_index('Date', inplace=True)

# Select features (You can add more relevant features if present in your dataset)
X = stocks_data.drop(['Close'], axis=1) # Features (Remove 'Close' which is the target)
y = stocks_data['Close'] # Target variable

# Perform one-hot encoding for categorical features
X = pd.get_dummies(X)

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Create and train the Random Forest Regressor
rf = RandomForestRegressor(n_estimators=100, random_state=42)
rf.fit(X_train, y_train)

# Make predictions
y_pred = rf.predict(X_test)

# Evaluate the model
mse = mean_squared_error(y_test, y_pred)
print(f'Mean Squared Error: {mse}')

# Plot the actual vs predicted prices
plt.figure(figsize=(10, 6))
plt.plot(y_test.values, label='Actual Prices')
plt.plot(y_pred, label='Predicted Prices')
plt.legend()
plt.show()
```



Mean Squared Error: 1.1492446760200026

